

Ranges, Partitioning, and Limitations

PgDay Austin 2016

Corey Huinker

What is this talk about?

An overview of what Range Types are and what they can do.

A series of gripes about what they can't do.

Cool uses for Range Types in my work at Moat (<http://moat.com>).

Why are Range Types Important?

- They allow your data to more accurately convey meaning.
- They allow your code to more accurately convey intention.
- Indexability, Exclusion constraints
- No other RDBMS has them ^[1], giving PostgreSQL an expressive advantage.

[1] - I haven't looked too hard.

Range Basics: Bounds

Ranges behave like and are denoted by standard mathematical Interval Notation.

Notation	Means	Notation	Means
$(x$	values $> x$	$[x$	values $\geq x$
$y)$	values $< y$	$y]$	values $\leq y$
$(,$	No lower bound	$,)$	No upper bound
$(,)$	everything	empty	No values


Constructing Ranges

Casting from text:

```
select '[low,high] '::rangetype
```

```
select '[low,)' '::rangetype
```

Omitting a bound means
unbounded, regardless of inc/excl




Creation through constructor function

```
select rangetype(low,high,'[]')
```

```
select rangetype(null,high,'[]')
```

Nulling a bound is the same as
omitting it.



Note: no polymorphic constructor

```
select to_range(null::rangetype,low,high,'[]');
```

NOPE



Range Basics: Existing Types

- `int4range`: Range of integer
- `int8range`: Range of bigint
- `numrange`: Range of numeric
- `tsrange`: Range of timestamp without time zone
- `tstzrange`: Range of timestamp with time zone
- `daterange`: Range of date
- ~~`boolrange`~~: Range of boolean
- ~~`textrange`~~: range of text

Why no textrange type?

- Collation Sequences.
 - Would need one textrange per collation sequence.
- No telling how many collations are installed.
 - Or what order they were installed in.
- Need one oid per range type, just like any other type.
- Would have to pre-allocate them with static type definitions.
- Not going to burn that many oids on a bunch of maybes.
 - So just define one type per collation sequence that you'll need
 - You probably only need "C" and maybe one other.

```
create type textrange_c as range (subtype = text, collation = "C");
```

Attribute functions:

Ranges can be decomposed into their component attributes.

```
# create temp table temps(state text, rng numrange);
# insert into temps values ('ice',          '(',32.0)'),
                           ('water',        numrange(32.0,212.0,'[])'),
                           ('steam',        numrange(212.0,null)),
                           ('heat death', 'empty');
```

```
# select * from temps;
```

state	rng
ice	(,32.0)
water	[32.0,212.0)
steam	[212.0,)
heat death	empty

Attribute functions In Action:

```
# \pset null '⌘'
```

```
Null display is "⌘".
```

This is really useful when sharing examples, but might be confusing if you think that's a currency symbol.

```
# select state, lower(rng) as low, lower_inc(rng) as low_inc,  
lower_inf(rng) as low_inf, upper(rng), upper_inc(rng), upper_inf(rng),  
isempty(rng) as empty from temps;
```

state	low	low_inc	low_inf	upper	upper_inc	upper_inf	empty
ice	⌘	f	t	32.0	f	f	f
water	32.0	t	f	212.0	f	f	f
steam	212.0	t	f	⌘	f	t	f
heat death	⌘	f	f	⌘	f	f	t

(4 rows)

Operators: =, <>

Discrete ranges normalize to the [] bound via the defined *canonical* function, and are then tested for equivalence. Continuous ranges do not have a *canonical* function, and are tested as-is.

Expression	Result
<code>select '(1,10]':::int4range;</code>	<code>[2,11)</code>
<code>select '[yesterday,today]':::daterange = '[yesterday,tomorrow)':::daterange;</code>	<code>t</code>
<code>select '[1,3]':::numrange = '[1,4)':::numrange;</code>	<code>f</code>
<code>select '[1,3]':::numrange = '[1,3.00000000000000000000001)':::numrange;</code>	<code>f</code>

Operators: <, <=, >, >=

- Test lower bound scalar first, then use upper bound as a tiebreaker
 - Which isn't really intuitive, but then again neither are the alternatives:
 - Median?
 - Number of (discrete) values contained?
- Therefore, not generally useful for userland queries.
- Used internally for indexing.

Operator <<

- "Strictly to the left of"
- $a << b$ if normalized upper bound of a is $<$ normalized lower bound of b

```
# select '[1,3) '::int4range << '[3,5) '::int4range as a1,  
        '[1,3] '::int4range << '[3,5) '::int4range as a2;
```

a1		a2
-----	+	-----
t		f

Operator >>

- "Strictly to the right of"
- $a \gg b$ if normalized lower bound of a is $>$ normalized upper bound of b

```
# select '[today,tomorrow) '::daterange >>
        '[yesterday,today) '::daterange as a1,
        '[today,tomorrow) '::daterange >>
        '[yesterday,today] '::daterange as a2;
```

a1 | a2

-----+-----

t | f

Operator <

- "Does not extend to the right of"
- No element of a is $>$ greatest element of b

```
# select daterange('[today,tomorrow)') <
          daterange('[yesterday,today)') as x,
          int4range('[10,20)') < int4range('[10,20]') as y;
```

```
x | y
---+---
f | t
(1 row)
```

Operator &>

- "Does not extend to the left of"
- No element of a is $<$ least element of b

```
# select '[3,10)>::int4range &> '[1,4)>::int4range as x,  
        '[0,10)>::int4range &> '[1,4)>::int4range as y;
```

```
x | y
```

```
---+---
```

```
t | f
```

Operator – | –

- "adjacent"
- There is no overlap nor space between *a* and *b*.
- It doesn't matter which range is lower

```
# select '[4,10) '::int4range -|- '[1,4) '::int4range as x,  
        '[1,3] '::int4range -|- '[5,10] '::int4range as y,  
        '[1,10] '::int4range -|- '[5,15] '::int4range as z;
```

```
x | y | z
```

```
---+---+---
```

```
t | f | f
```


Operators <@ and @>

- "contains", same as the geometric operators
- The value or range on the pointy side fits entirely within the range on the @ side
- It doesn't matter which range is lower

```
# select 1 <@ '[1,4]':::int4range as u,  
        '[20,30)':::int4range <@ '[1,100]':::int4range as v,  
        'infinity':::date <@ '(',')':::daterange as w,  
        '(',')':::int4range @> 'empty':::int4range as x,  
        '(',')':::int4range @> null as y;
```

u		v		w		x		y
---	+	---	+	---	+	---	+	---
t		t		t		t		⊗

Operator &&

- "overlap", same as the geometric operator
- At least one value can fit in both ranges

```
# select '[20,30) '::int4range && '[1,100] '::int4range as v,  
        '(',') '::int4range      && 'empty'::int4range as x;
```

v		x
---+---		
t		f

```
# select 'empty'::int4range <@ '(',') '::int4range as v,  
        'empty'::int4range && '(',') '::int4range as x;
```

v		x
---+---		
t		f

Operator + (and the range_merge() function)

- Union: All elements in both, if there are no gaps

```
# select int4range(1,4) + int4range(2,10) as x;
```

x

[1,10)

```
# select int4range(1,2) + int4range(99,100) as y;
```

ERROR: result of range union would not be contiguous

```
# select range_merge(int4range(1,2),int4range(99,100)) as z;
```

z

[1,100)



New in 9.5!

Available for earlier version in range_type_functions on PGXN

Operator *

- Intersection: all elements in common, if any

```
# select int4range(1,4) * int4range(4,100) as x,  
        int4range(1,4,'[]') * int4range(4,100) as y;
```

```
      x      |      y  
-----+-----  
empty | [4,5)
```

Operator –

- Difference: all elements in *a* but not in *b*
- Will raise an error if the difference would return 2 disjoint sets

```
# select int4range(1,100) - int4range(1,10) as x;
```

```
      x
```

```
-----
```

```
[10,100)
```

```
# select int4range(1,100) - int4range(2,10) as x;
```

```
ERROR:  result of range difference would not be contiguous
```

Missing Function: `range_split()`

- Same as the `-` operator, but returning the left side remainder and right side remainder
- returns an array of the resulting ranges
- a SRF would be nice too.

```
hypothetical# select range_split('[1,100]':::int4range,  
                                '[2,4]':::int4range) as x;
```

x

```
-----  
{ [1,2), [2,5), [5,100] }
```

Missing Operators =|, |=

Operators to test whether two ranges share a lower (=|) bound or upper bound (|=)

```
hypothetical# select '[1,4]':::int4range =| '[1,10]':::int4range as w,  
                    '[1,4]':::int4range =| '(1,10]':::int4range as x,  
                    '[1,4]':::int4range |= '(',4)':::int4range as y,  
                    '[1,4]':::int4range |= '(',4)':::int4range as z;
```

w	x	y	z
---	+	---	+
t	f	t	f

Missing Operators: elem <<, >>

- Same as the current <</>> operators, but allow the one arg to be a scalar.
- May be a problem for existing bitshift operators

```
hypothetical# select 1::integer << '[1,10]':::int4range as w,  
                    1::integer << '(1,4]':::int4range as x,  
                    4::integer >> '[1,4]':::int4range as y,  
                    4::integer >> '(,4)':::int4range as z;
```

```
 w | x | y | z  
---+---+---+---  
 f | t | f | t
```

Can be simulated by creating a singleton range:
`int4range(1,1,'[]') << int4range(2,11,'[]')`

Missing Operator: `elem <=> range`

- Returns 0 if element a `<@` range b .
- -1 if a `<<` b , 1 if a `>>` b
- basically `strcmp()` but for ranges

```
hypothetical# select 1::integer <=> '[1,10]':::int4range as w,  
                    1::integer <=> '(1,4]':::int4range as x,  
                    4::integer <=> '[1,4]':::int4range as y,  
                    4::integer <=> '(,4)':::int4range as z;
```

w	x	y	z
0	-1	0	1

Implemented as
`element_range_comp()` in
`range_type_functions` on PGXN

Missing Functions: `is_singleton()`

- Return true if the range can contain only one element.

```
# select is_singleton('[4,5)::int4range);  
is_singleton  
-----  
t
```

```
# select is_singleton('[4,5)::int4range);  
is_singleton  
-----  
f
```

Found in
`range_type_functions` on
PGXN

Missing Functions: get bounds

- Represent either or both bounds conditions as SQL
- Helpful when constructing `CHECK` / `WHERE` clauses or dealing with foreign systems that don't support that range type or ranges in general.

```
# with t(c) as (values('[1,4]':::int4range))
  select get_lower_bound_condition_expr(c) as l,
         get_upper_bound_condition_expr(c) as u,
         get_bounds_condition_expr(c, 'zz') as b from t;
```

Found in
`range_type_functions` on
PGXN

l		u		b
-----+-----+-----				
<code>x >= '1':::integer x < '5':::integer zz >= '1':::integer and zz < '5':::integer</code>				

Partitioning by Ranges Use Case

Use case is a series of "typeahead search" tables:

- Hundreds of millions of rows.
- Grouped by a taxonomy of 5 text strings of increasing length.
- The searchable text is usually 5-20 words per record
- Need a way to partition the table, but only text types available.
- Distribution is highly uneven along strict alphabetical lines.

Text Range Partitioning Advantages

- partitions have smaller GIN indexes on the searchable columns, so smaller `BitmapAnd` steps
- Ability to isolate very large clients.
- Search dataset evolves over time the lumps in the data move, but slowly.
- Partition maintenance only when data is starting to skew, much different from timeseries.

[illegible]

range_partitioning module

- On PGXN
- Functions closely match those in pg_partman.
 - `create_parent(table, column_name)`
 - starts with implied range of (,)
 - `create_partition(table, new_range)`
 - new partition range must be perfect subset of an existing range, and match lower or upper bound.
 - `drop_partition(lost_part, kept_part)`
 - merge all data from lost_part into kept_part

range_partitioning module

- `SELECT / INSERT / UPDATE` queries are transparent.
- Does trigger function for transparent `INSERT`
- Probably better having bulk loads separated by partitioned value, and probing for the destination partition with `get_destination_partition()`, if possible.
- The `create_parent()` function cannot seamlessly derive the base type if more than one range type has that base type.
- Ranges are specified as un-casted text strings.

range_partitioning example

Use case: Message board for fans of TV shows. The site's users skew heavily towards certain niche shows.^[1]

```
/* Turn existing table into a parent table. One partition with range (,) */  
select range_partitioning.create_parent('public.spoiler_alerts',  
                                       'tv_show_name');
```

```
/* Create a partition just for the show ARCHER, but all new partitions must  
share an edge with an existing partition, so you may need to explicitly  
create more than one */
```

```
select range_partitioning.create_partition('public.spoiler_alerts',  
                                           '(,ARCHER)');  
  
select range_partitioning.create_partition('public.spoiler_alerts',  
                                           '[ARCHER,ARCHER]');
```

[1] The niche is defined as "Shows I can name".

range_partitioning example (part 2)

[illegible]

```
/* Create a partition just for the show RICK_AND_MORTY, again sharing an
edge */
```

[illegible]

range_partitioning: partition list

```
# select partition_number, range
  from range_partitioning.partition
  where master_class = 'public.spoiler_alerts'::regclass;
```

partition_number	range
0	(GAME_OF_THRONES,RICK_AND_MORTY)
1	(,ARCHER)
2	[ARCHER,ARCHER]
3	(ARCHER,GAME_OF_THRONES]
4	(RICK_AND_MORTY,)
5	[RICK_AND_MORTY,RICK_AND_MORTY]

range_partitioning type discovery

The `create_parent(table, column)` function doesn't need to have the range type specified **if** only one range type would work for that column.

```
/* if this returns more than one row, then we have to specify a range type
*/
select  rt.rngtypeid
from    pg_attribute a
join    pg_range rt
on      rt.rngsubtype = a.atttypeid
and     rt.rngcollation = a.attcollation
where   a.attrelid = 'my_schema.my_parent_table'::regclass
and     a.attname = 'my_partitioning_column';
```

Complex Range Partitioning

- Possible to partition on ranges of complex types
 - That complex type must exist in the table itself, it can't be more than one column
 - So re-expose the components in a view.

```
# create type quite_complex as (a text collate "C", b text collate "C",  
                                c text collate "C", d text collate "C");  
  
CREATE TYPE  
# create type qc_range as range (subtype = quite_complex);  
CREATE TYPE  
# select  
'["(Abel,Baker,Charlie,Delta''s)","(Walter,X-Ray,Yellow,)")'::qc_range;  
      qc_range  
-----  
["(Abel,Baker,Charlie,Delta's)","(Walter,X-Ray,Yellow,)")
```

Future Direction: range_partitioning

- Add functions to predict proper partition ranges for equal-ish row counts
 - `width_buckets()` works ok, but will sometimes skip some buckets entirely. You ask for 16 partitions, get ~13.
- Add functions to analyze existing partitions for skew
- Become obsolete.
 - Native Partitioning coming to PostgreSQL in v10, probably.
 - Existing work supports ranges but not range syntax.
 - Using BOUNDS syntax, not ranges.
 - Might only support '[' ranges.

Links

Range Partitioning extension:

PGXN: http://pgxn.org/dist/range_partitioning/

GitHub: https://github.com/moat/range_partitioning

Range Type Functions:

PGXN: http://pgxn.org/dist/range_type_functions/

GitHub: https://github.com/moat/range_type_functions [*]

[*] - likely moving to new ownership